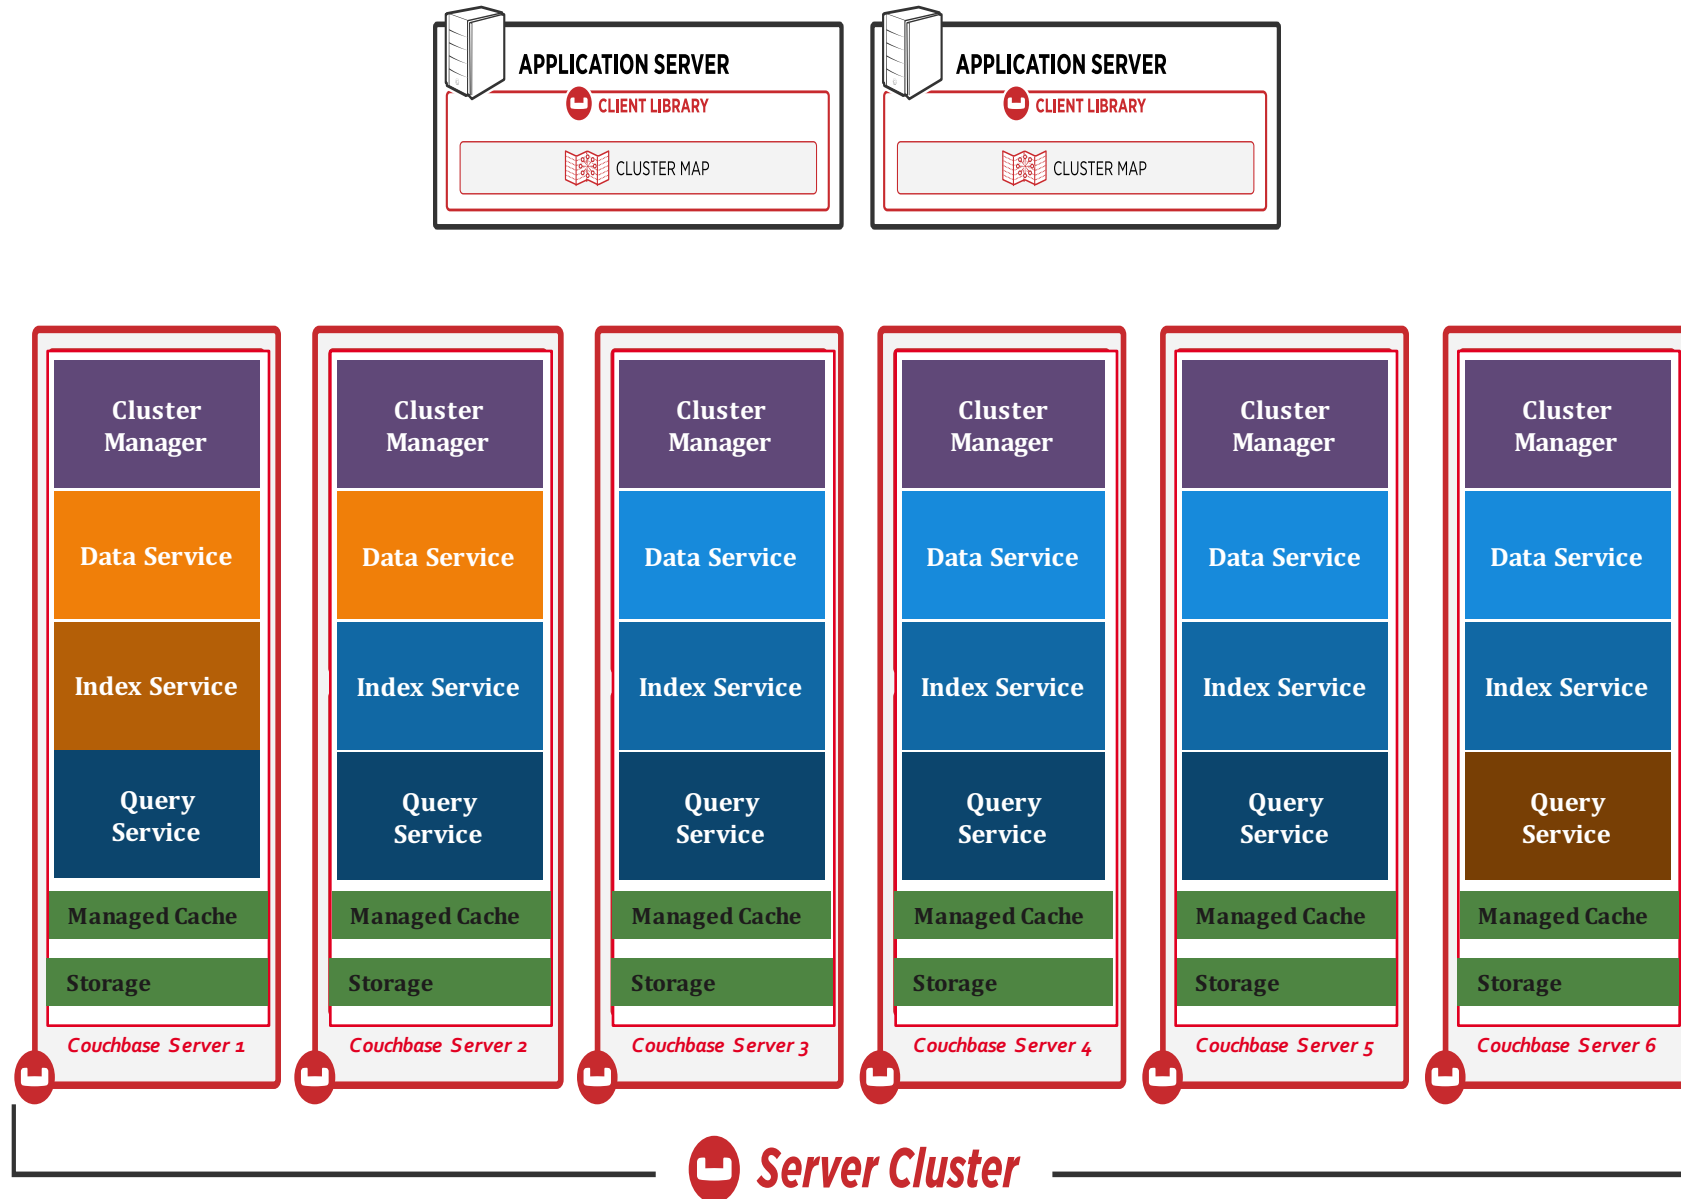


# **Global Secondary Indexes new high performance indexer**

# Couchbase Server Cluster Architecture



- **Multiple Indexers**

- **GSI – Index Service**

New indexing for N1QL for low latency queries without compromising on mutation performance (insert/update/delete)



Independently partitioned and independently scalable indexes in Indexing Service

- **Map/Reduce Views – Data Service**

Powerful programmable indexer for complex reporting and indexing logic.

Full partition alignment and paired scalability with Data Service.

- **Spatial View – Data Service**

Incremental R-tree indexing for powerful bounding-box queries

Full partition alignment and paired scalability with Data Service

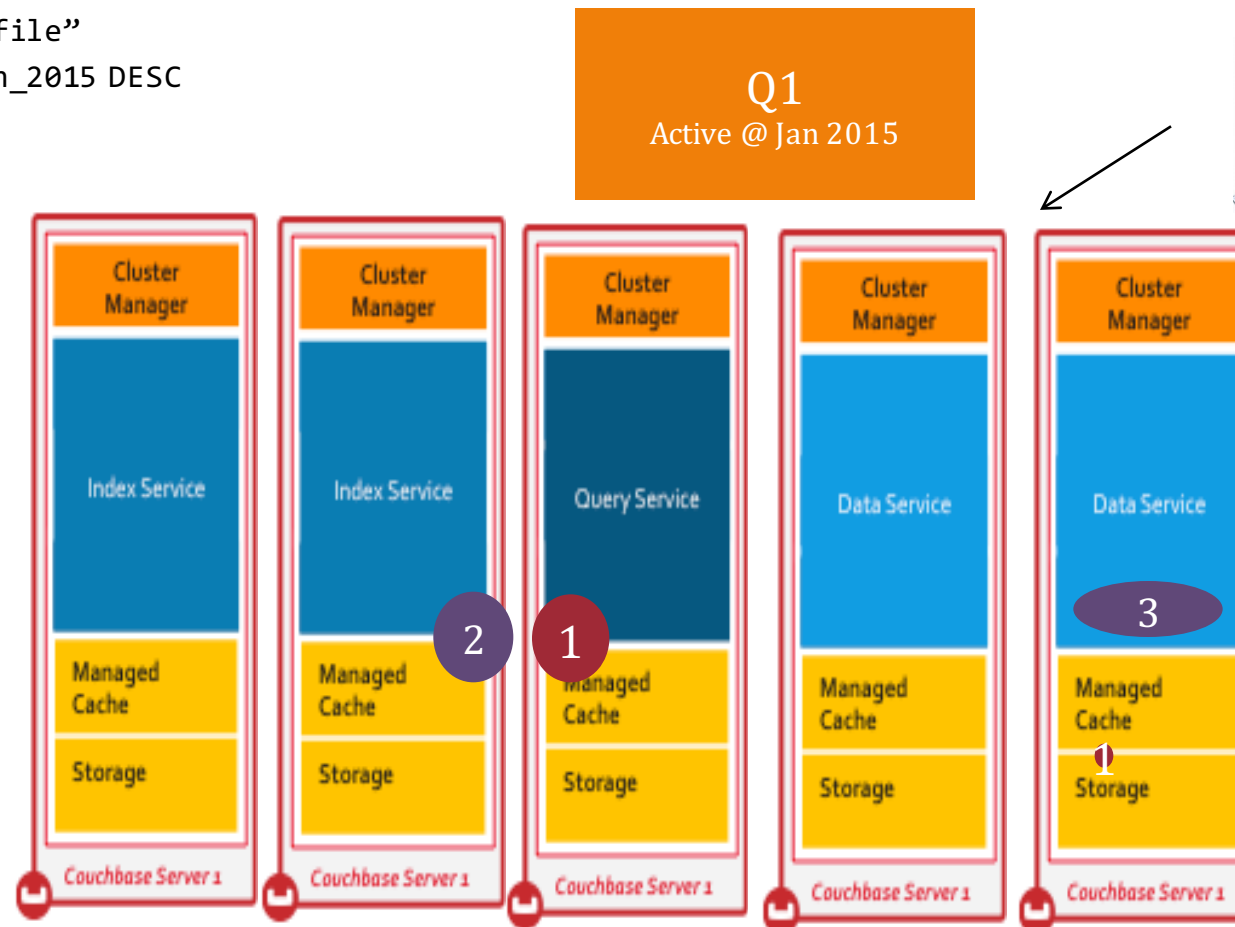
# Query and Index with GSI

```
Create INDEX ON Customer_bucket(customer_name, total_logins.jan_2015)
WHERE type="customer_profile";
```

```
SELECT customer_name, total_logins.jan_2015
FROM customer_bucket
WHERE type="customer_profile"
ORDER BY total_logins.jan_2015 DESC
LIMIT 10;
```

Q1: Execution Plan on N nodes

- Execute Q1 on N1QL Service node
- Scan index on Index Service node



## What are Global Secondary Indexes?

*High performance indexes for low latency queries with powerful caching, storage and independent placement.*

- Power of GSI
  - **Fully integrated into N1QL** Query Optimization and Execution
  - **Independent Index Distribution** for Limiting scatter-gather
  - **Independent Scalability** with Index Service
  - **Powerful caching and storage** with ForestDB

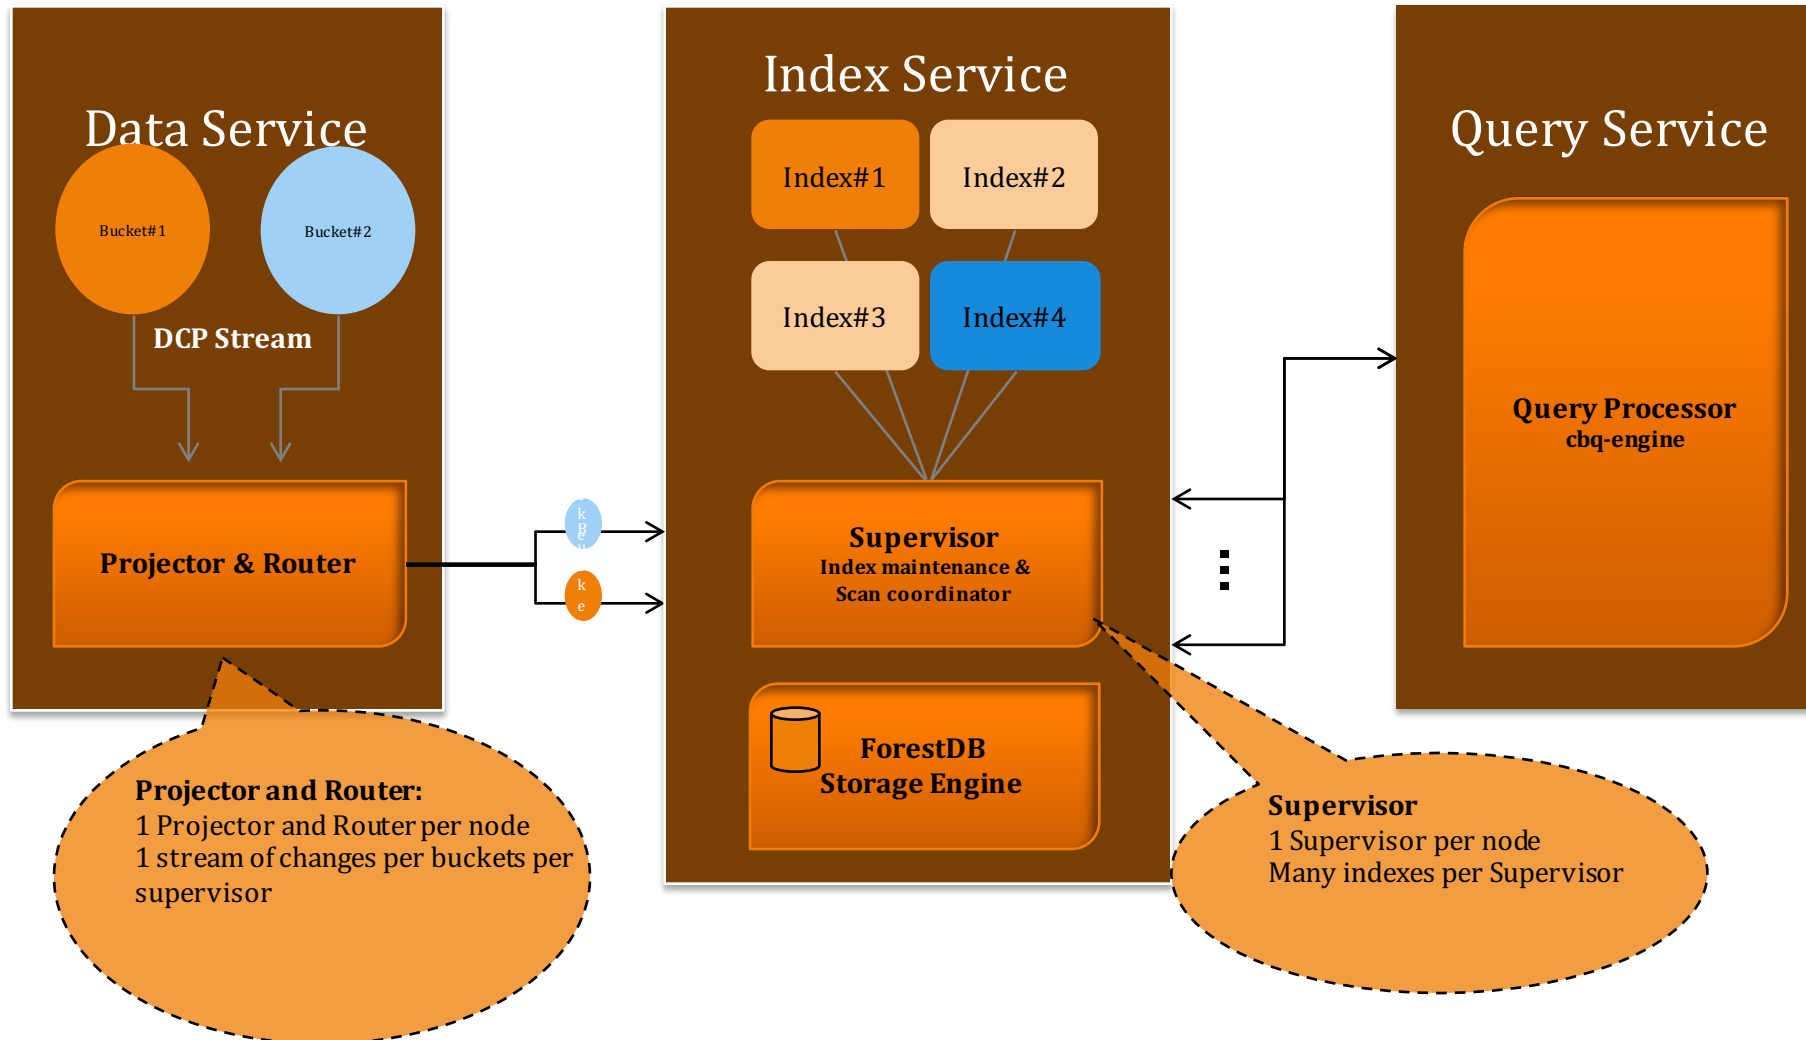
# Which to choose – GSI vs Views

| Workloads             | New GSI in v4.0                                | Map/Reduce Views        |
|-----------------------|--|-------------------------|
| Complex Reporting     | Just In Time Aggregation                       | Pre-aggregated          |
| Workload Optimization | Optimized for Scan Latency & Throughput        | Optimized for Insertion |
| Flexible Index Logic  | N1QL Functions                                 | Javascript              |
| Secondary Lookups     | Single Node Lookup                             | Scatter-Gather          |
| Tunable Consistency   | Staleness false or ok or everything in between | Staleness false or ok   |

| Capabilities         | New GSI in v4.0                          | Map/Reduce Views               |
|----------------------|--|--------------------------------|
| Partitioning Model   | Independent – Indexing Service           | Aligned to Data – Data Service |
| Scale Model          | Independently Scale Index Service        | Scale with Data Service        |
| Fetch with Index Key | Single Node                              | Scatter-Gather                 |
| Range Scan           | Single Node                              | Scatter-Gather                 |
| Grouping, Aggregates | With N1QL                                | Built-in with Views API        |
| Caching              | Managed                                  | Not Managed                    |
| Storage              | ForestDB                                 | Couchstore                     |
| Availability         | Multiple Identical Indexes load balanced | Replica Based                  |

# GSI Architecture





# GSI Lifecycle

# Indexing Lifecycle

- Primary vs Secondary
  - Primary Index is a full list of document keys within a given bucket

```
CREATE PRIMARY INDEX index_name
ON bucket_name
USING GSI|VIEW
WITH `{"nodes": ["node_name"], "defer_build":true}`; //GSI-ONLY
```

- Secondary Index is an index on a field/expression on a subset of documents for lookups

```
CREATE INDEX index_name
ON bucket_name (field/expression, ...)
USING GSI|VIEW
WHERE filter_expressions
WITH `{"nodes": ["node_name"], "defer_build":true}`; //GSI-ONLY
```

- Index building can be deferred to build multiple indexes all at once with greater scan efficiency.

```
CREATE INDEX ... WITH {..."defer_build":true};  
CREATE INDEX ... WITH {..."defer_build":true};  
...  
BUILD INDEX ON bucket_name(index_name1, ...) USING GSI;
```

# **GSI Partitioning and Placement**

# GSI Placement and Partitioning


## Place GSI Indexes using NODES clause

- *Each GSI reside on 1 node*
  - You can specify the node using nodes clause

```
CREATE INDEX i1 ... WITH {"nodes": "node1"};
```

- You can scale out the index by creating identical indexes (load balanced)

```
CREATE INDEX i1 ... WITH {"nodes": "node1"};  
CREATE INDEX i2 ... WITH {"nodes": "node2"};
```



## Partition Indexes Manually with WHERE clause

- You can partition with the WHERE clause and place on various nodes for scaling out

```
CREATE INDEX i1 ... WHERE zipcode between "94000" and "94999" ...  
CREATE INDEX i2 ... WHERE zipcode between "95000" and "96000" ...
```

# **GSI Availability and Rebalance**



## Use multiple identical indexes for availability

- GSI Availability
  - Create multiple identical indexes on separate nodes for availability
  - GSI will auto divert traffic if any copy goes down
- GSI & Rebalance
  - Removing/Failing a node with index service, remove the GSI indexes on that node
  - Adding a node with index service, won't automatically move some indexes to the node.

# Monitoring GSI

# Monitoring GSI Indexes

- Index Size and Maintenance Stats
- Index Scan Stats

[tos.couchbase](#) > [Servers](#) > Statistics

## ▼ Index Stats: Index\_tos\_beer\_city



0 items scanned/sec  
per server



28.8MB disk size  
per server



209KB data size  
per server



2.78KB memory used  
per server



0 total mutations remaining  
per server



0 drain rate items/sec  
per server



1412 total indexed items  
per server



152B average item size  
per server



99.2 % fragmentation  
per server



0 requests/sec  
per server



0 bytes returned/sec  
per server



0 avg scan latency(ns)  
per server



0 cache resident percent  
per server



0 index cache miss ratio  
per server

- Indexing Complex Types
  - Sub-documents attributes

```
CREATE INDEX ifriend_id ON default(friends.id)  
USING GSI;
```

```
SELECT * FROM default  
WHERE friends.id= "002819";
```

```
{  
  "id": "0000000000000001",  
  "desc": "---",  
  "type": "friends",  
  "tags": [0,1,2,3,4,5,6,7,8,9],  
  "friends": {  
    "id": "002819",  
    "class": "005"  
  }  
}
```

- Indexing Complex Types
  - Compound Keys

```
CREATE INDEX ifriends_id_class ON  
default(friends.id, friends.class) USING GSI;  
  
SELECT * FROM default  
WHERE friends.id="002819" and friends.class="005";
```

```
{  
  "id": "0000000000000001",  
  "desc": "---",  
  "type": "friends",  
  "tags": [0,1,2,3,4,5,6,7,8,9],  
  "friends": {  
    "id": "002819",  
    "class": "005"  
  }  
}
```

- Indexing Complex Types
  - Sub-documents

```
CREATE INDEX ifriend ON default(friends) USING GSI;
```

```
SELECT * FROM default  
WHERE friends= {"class": "005","id":"002819"};
```

```
{  
  "id": "0000000000000001",  
  "desc": "---",  
  "type": "friends",  
  "tags": [0,1,2,3,4,5,6,7,8,9],  
  "friends": {  
    "id": "002819",  
    "class": "005"  
  }  
}
```

# Types of Indexes

- CREATE PRIMARY INDEX ON `travel-sample`;

## Metadata for Primary Index

```
SELECT * FROM system:indexes WHERE name = '#primary';
```

Results:

```
[
  {
    "indexes": {
      "datastore_id": "http://127.0.0.1:8091",
      "id": "1b7ac1abf01d9038",
      "index_key": [],
      "is_primary": true,
      "keyspace_id": "travel-sample",
      "name": "#primary",
      "namespace_id": "default",
      "state": "online",
      "using": "gsi"
    }
  }
]
```



# Named Primary Index

- CREATE PRIMARY INDEX def\_primary ON `travel-sample`;

```
SELECT meta().id AS documentkey, `travel-sample` airline
FROM `travel-sample`
WHERE type = 'airline'
LIMIT 1;
```

Results:

```
[
  {
    "airline": {
      "callsign": "MILE-AIR",
      "country": "United States",
      "iata": "Q5",
      "icao": "MLA",
      "id": 10,
      "name": "40-Mile Air",
      "type": "airline"
    },
    "documentkey": "airline_10"
  }
]
```

- `CREATE INDEX travel_name ON `travel-sample`(name);`

```
CREATE INDEX travel_geo on `travel-sample`(geo);

# get is an object embedded within the document such as:
# "geo": {
#   "alt": 12,
#   "lat": 50.962097,
#   "lon": 1.954764
# }
```

- `CREATE INDEX travel_info ON `travel-sample`(name,type,id,icoo,iata);`

Each of the keys can be a simple scalar field, object, or an array. For the index filtering to be exploited, the filters have to use respective object type in the query filter.

- CREATE INDEX travel\_cxname ON `travel-sample`(LOWER(name));

```
{
  "#operator": "Parallel",
  "~child": {
    "#operator": "Sequence",
    "~children": [
      {
        "#operator": "Filter",
        "condition": "(lower(`travel-sample`.`name`)) = \"john\""
      },
      {
        "#operator": "InitialProject",
        "result_terms": [
          {
            "expr": "self",
            "star": true
          }
        ]
      }
    ]
  }
},
```

- CREATE INDEX travel\_sched ON `travel-sample`  
(ALL DISTINCT ARRAY v.day FOR v IN schedule END);

```
schedule:
[
  {
    "day" : 0,
    "special_flights" :
    [
      {
        "flight" : "AI111", "utc" : "1:11:11"
      },
      {
        "flight" : "AI222", "utc" : "2:22:22"
      }
    ]
  },
  {
    "day": 1,
    "flight": "AF552",
    "utc": "14:41:00"
  }
]
```

- CREATE INDEX travel\_info ON `travel-sample`(name, id, icao, iata)  
WHERE type='airline';

```
{
  "airline": {
    "callsign": "MILE-AIR",
    "country": "United States",
    "iata": "Q5",
    "icao": "MLA",
    "id": 10,
    "name": "40-Mile Air",
    "type": "airline"
  },
  "documentkey": "airline_10"
}
```

- `CREATE INDEX i1 ON `travel-sample`(LOWER(name),id, icoo)  
WHERE type = 'airline';`
- `CREATE INDEX i2 ON `travel-sample`(LOWER(name),id, icoo)  
WHERE type = 'airline';`
- `CREATE INDEX i3 ON `travel-sample`(LOWER(name),id, icoo)  
WHERE type = 'airline';`

All three indexes have identical keys and an identical WHERE clause; the only difference is the name of these indexes. You can choose their physical location using the WITH clause of the CREATE INDEX statement.

providing scale-out, multi-dimensional scaling, performance, and high availability

- Index selection for a query solely depends on the filters in the WHERE clause of your query.
- After the index selection is made, the query engine analyzes the query to see if it can be answered using only the data in the index
- `CREATE INDEX travel_info ON `travel-sample`(name, id, icoo, iata)`



# Memory-Optimized Index Storage

- Memory-optimized index-storage allows high-speed maintenance and scanning; since the index is kept fully in memory at all times.
- A snapshot of the index is maintained on disk, to permit rapid recovery if node-failures are experienced.
- less suitable for nodes where memory is constrained;

# Standard Index Storage

- *Standard* is the default storage-setting for Secondary Indexes: the indexes are saved on disk; in a disk-optimized format that uses both memory and disk for index-update and scanning.
- The performance of standard index storage depends on overall I/O performance.

Settings are established at cluster-initialization for all indexes on the cluster, across all buckets. Following cluster-initialization, to change from one setting to the other, all nodes running the Index Service must be removed. If the cluster is single-node, uninstall and reinstall Couchbase Server.

## Index Storage Mode

☐ Standard Global Secondary

☒ Memory-Optimized ⓘ

Changing the optimization mode of global indexes is not supported when index service nodes are present in the cluster. Please remove all index service nodes to change this option.

# **Labs:**

## **Query Workbench & Index– 60 Minutes(D)**